

Demonstrate partitioning by inserting data with different partition keys
Cassandra Exp.4

Demonstrating Partitioning in Cassandra

1. Create a keyspace
2. Create a table with a partition key
3. Insert records with different partition keys
4. Observe how data groups by partitions
5. Query data by partition key

1. Create Keyspace

```
CREATE KEYSPACE IF NOT EXISTS partition_demo  
WITH REPLICATION = {  
    'class': 'SimpleStrategy',  
    'replication_factor': 1  
};  
  
USE partition_demo;  
  
cqlsh> CREATE KEYSPACE IF NOT EXISTS partition_demo  
... WITH REPLICATION = {  
...     'class': 'SimpleStrategy',  
...     'replication_factor': 1  
... };  
cqlsh> USE partition_demo;
```

2. Create a Table with a Clear Partition Key

The table below uses **department** as the partition key.
All students in the same department will be stored in the same partition.

```
CREATE TABLE IF NOT EXISTS students_by_department (  
    department text,  
    student_id UUID,  
    name text,  
    year int,  
    PRIMARY KEY (department, student_id)  
);
```

```
cqlsh:partition_demo> CREATE TABLE IF NOT EXISTS students_by_department (
...     department text,
...     student_id UUID,
...     name text,
...     year int,
...     PRIMARY KEY (department, student_id)
... );
```

Explanation:

- **Partition Key:** department
- **Clustering Key:** student_id

Thus:

- A partition is created per department.
- Each partition stores multiple students belonging to that department.

3. Insert Data with Different Partition Keys

Department: Computer Science (Partition 1)

```
INSERT INTO students_by_department (department, student_id, name, year)
VALUES ('CS', uuid(), 'Arjun', 3);

cqlsh:partition_demo> INSERT INTO students_by_department (department, student_id, name, year)
... VALUES ('CS', uuid(), 'Arjun', 3);

INSERT INTO students_by_department (department, student_id, name, year)
VALUES ('CS', uuid(), 'Meghana', 2);

cqlsh:partition_demo> INSERT INTO students_by_department (department, student_id, name, year)
... VALUES ('CS', uuid(), 'Meghana', 2);
```

Department: Electronics (Partition 2)

```
INSERT INTO students_by_department (department, student_id, name, year)
VALUES ('ECE', uuid(), 'Shinchana', 1);

cqlsh:partition_demo> INSERT INTO students_by_department (department, student_id, name, year)
... VALUES ('ECE', uuid(), 'Shinchana', 1);

INSERT INTO students_by_department (department, student_id, name, year)
VALUES ('ECE', uuid(), 'Rohit', 3);

cqlsh:partition_demo> INSERT INTO students_by_department (department, student_id, name, year)
... VALUES ('ECE', uuid(), 'Rohit', 3);
```

Department: Mechanical (Partition 3)

```
INSERT INTO students_by_department (department, student_id, name, year)
VALUES ('ME', uuid(), 'Vishal', 4);

cqlsh:partition_demo> INSERT INTO students_by_department (department, student_id, name, year)
... VALUES ('ME', uuid(), 'Vishal', 4);
```

```

INSERT INTO students_by_department (department, student_id, name, year)
VALUES ('ME', uuid(), 'Sahana', 2);

cqlsh:partition_demo> INSERT INTO students_by_department (department, student_id, name, year)
... VALUES ('ME', uuid(), 'Sahana', 2);

```

These inserts create **three separate partitions**:

- Partition for key **CS**
- Partition for key **ECE**
- Partition for key **ME**

Each partition holds only its department's rows.

4. Query Data by Partition Key

Retrieve only CS students

```
SELECT * FROM students_by_department
```

```
WHERE department = 'CS';
```

```
cqlsh:partition_demo> SELECT * FROM students_by_department
... WHERE department = 'CS';
```

department	student_id	name	year
CS	e13ad99b-c559-40ef-aad5-ec9b4a82797e	Arjun	3
CS	f751db68-0d72-4197-90fc-c6a202300419	Meghana	2

Retrieve only ECE students

```
SELECT * FROM students_by_department
```

```
WHERE department = 'ECE';
```

```
cqlsh:partition_demo> SELECT * FROM students_by_department
... WHERE department = 'ECE';
```

department	student_id	name	year
ECE	6d76831e-ff78-43e3-9999-84b175eac5c8	Rohit	3
ECE	ea3731a7-b108-4066-b91f-746766ec95fd	Shinchana	1

Retrieve only ME students

```
SELECT * FROM students_by_department
```

```
WHERE department = 'ME';
```

```
cqlsh:partition_demo> SELECT * FROM students_by_department  
        ... WHERE department = 'ME';
```

department	student_id	name	year
ME	2c0d85a5-2d6f-4165-b3c8-faf4c0d31a20	Vishal	4
ME	bbb19c89-3340-4cf8-9ae8-1b2753ce8e68	Sahana	2

Each query reads from a **single partition**, demonstrating how Cassandra organizes data based on the partition key.

5. Understanding the Partitioning Behavior

1. Partition Key Determines Location

Cassandra hashes the partition key:

hash(department)

This decides which node stores that partition.

2. All Rows in the Same Partition Are Stored Together

Example:

All students in ECE stay in one physical partition on a node.

3. Efficient Reads

Querying by partition key is efficient because Cassandra directly finds the partition without scanning the whole table.

4. Inserts Scale Automatically

Adding a new department automatically creates a new partition.

6. Verification Using System Tables (Optional)

List partitions for the table

```
SELECT * FROM system.schema.tables
```

```
WHERE keyspace name='partition demo';
```

```
cqlsh:partition_demo> SELECT * FROM system_schema.tables  
... WHERE keyspace_name='partition_demo';
```

```

keyspace_name | table_name           | additional_write_policy | allow_auto_snapshot | bloom_filter_fp_chance | caching
               | cdc | comment | compaction

               | compression

| crc_check_chance | dclocal_read_repair_chance | default_time_to_live | extensions | flags      | gc_grace_seconds | id
               | incremental_backups | max_index_interval | memtable | memtable_flush_period_in_ms | min_index_interval | read_repair | read_repair_chance | speculative_retry

partition_demo | students_by_department |          99p |           null |        0.01 | {'keys': 'ALL',
', 'rows_per_partition': 'NONE'} | null |          | {'class': 'org.apache.cassandra.db.compaction.SizeTieredCompactionStrategy',
'max_threshold': '32', 'min_threshold': '4'} | {'chunk_length_in_kb': '16', 'class': 'org.apache.cassandra.io.compress.LZ4Compressor'} |          1 |          0 |
0 |          {} | {'compound'} |          864000 | 00f088d0-0739-11f1-9122-cdf73f26f65f |           null |

2048 |     null |          0 |
128 |    BLOCKING |          0 |          99p

```

```
(1 rows)
partition_demo | students_by_department |      name |      none |          0x6e616d65 |      regular |      -1 | t
ext
partition_demo | students_by_department | student_id |      asc | 0x73747564656e745f6964 | clustering |      0 | u
uid
partition_demo | students_by_department |      year |      none |          0x79656172 |      regular |      -1 |
int
```

Inspect partition columns

```
SELECT * FROM system_schema.columns
```

```
WHERE keyspace_name='partition_demo'
```

```
AND table_name='students_by_department';
```

This confirms that **department** is indeed the partition key.